




MEET THE HOUNDS

ANUBIS

ACR Calibration

Anubis Hound · Category: Correlation

“Make risk scores reflect your business reality.”

The weigher of worth — Anubis weighs each asset's true criticality, so every priority in the pack is honest.

THE PROBLEM

Out-of-the-box asset criticality rarely matches reality — the test box scores like a crown jewel while the actual payment server sits at a default rating. And because every downstream tool prioritizes against that score, one miscalibrated number quietly skews every list your team works from.

WHAT IT DOES

Anubis adjusts Asset Criticality Rating (ACR) per tag — set a value, or increment and decrement by N — each change carrying a business justification. The result is the shared **criticality currency** the entire pack prioritizes against: calibrate once, and every hound's ranking gets smarter.

KEY CAPABILITIES

- **Tag-driven ACR set / +N / -N** — calibrate a whole business tier at once: set production to its true criticality, down-rate the lab, raise the payment segment — no asset-by-asset clicking, no spreadsheet round-trips.
- **Business justification on every change** — the “why” is recorded with the number, so a year later the audit answer is in the log rather than in someone's memory.
- **Feeds every prioritizer** — attack-path ranking, crypto-migration order, AI risk, and identity criticality all consume the calibrated score. One adjustment re-prioritizes the entire pack, coherently.
- **Gated writes** — every ACR change is proposed, human-approved, and logged before it lands anywhere.

HOW IT WORKS

Anubis operates on Tenable's ACR through tag-scoped adjustments — typically against business tiers your producer agents have already tagged (production, crown-jewel, lab, OT). All writes are proposed, human-approved, and logged with their justification.

WHY IT'S DIFFERENT

- **One calibration re-prioritizes the whole pack at once** — Fenrir's attack paths, Heimdall's migration order, Pythia's AI risk, and Janus's identity ranking all shift together.
- **A currency, not a report** — ACR is the number every agent already speaks; Anubis just makes it true.
- **Justified, not arbitrary** — every adjustment records the business reason, so audits get answers instead of shrugs.

PROOF POINTS

- In a reference environment of 268 assets, Anubis's calibrated ACR is what defined the **67 crown-jewel targets** (ACR ≥ 7) that Fenrir's 63 attack paths ranked against — and the criticality weighting behind Heimdall's migration roadmap.

Illustrative results from a demo lab — not a guarantee. Calibration quality depends on the business-tier tagging it builds on; garbage tiers in, garbage priorities out.

WORKS BETTER WITH

Anubis consumes tags from every producer — Janus, Mimir, Pythia, Cerberus, Charon — and powers every correlator: **Fenrir**, **Heimdall**, **Pythia**, and **On the Scent** all rank against its calibrated score.

WHO IT'S FOR

Vulnerability management leads whose prioritization is only as good as the criticality under it; risk owners who know which assets actually matter; CISOs who want one defensible criticality story across every report.

CALL TO ACTION

Let Anubis weigh your estate — every priority downstream depends on the scales.