



MEET THE HOUNDS

FENRIR

Attack Path

Fenrir Hound · Category: Correlation

“See the chains an attacker would follow — break them first.”

The great wolf fated to break every chain — Fenrir maps the chains an attacker would follow so you can break them first.

THE PROBLEM

Vulnerability lists and identity issues live in separate silos — nobody can see that an internet-exploitable box sits one hop from a crown-jewel server. Teams drown in thousands of findings with no sense of which chain actually gets an attacker to the crown jewels.

WHAT IT DOES

Fenrir chains exploitability × identity × reachability into ranked, breakable attack paths. It correlates exploitable and weak-auth footholds, same-subnet lateral movement (with credential pivots), and high-value crown-jewel targets ($ACR \geq 7$) — then ranks every path by risk, so the first chain you break is the one that matters most.

KEY CAPABILITIES

- **Foothold detection** — flags every asset an attacker could enter through: CISA-KEV or critical vulnerabilities, SSH password authentication, and default or blank credentials. The difference between “vulnerable” and “a way in.”
- **Reachability via /24 subnet adjacency** — maps which footholds sit within reach of which targets, so a critical vuln on an isolated segment stops outranking a medium one sitting next to your crown jewels.
- **Credential-pivot boost** — paths where reused or weak credentials let an attacker hop between hosts rank higher, mirroring how real intrusions actually move laterally.
- **Path risk score** — combines target ACR, entry strength, pivot potential, and KEV status into one number per path, so triage starts from a defensible ranking instead of a debate.
- **3-stage path table** — every chain laid out entry → pivot → target: readable in one glance by an analyst, explainable in one slide to a CISO.
- **Gated “Attack Path” tags + CSV export** — the endpoints of ranked paths become human-approved Entry/Target tags and exportable worklists that route straight into remediation.

HOW IT WORKS

Fenrir runs against the local navi.db built from your Tenable data: the **vulns** table (KEV cross-references, severity, plugins 149334 and 41028) and the **assets** table (IP, ACR). All writes are proposed, human-approved, and logged.

WHY IT'S DIFFERENT

- Turns **four separate agents' signals into one prioritized list** — exploitability, credentials, criticality, and software risk compound instead of competing.
- **Honest about reachability**: paths are inferred from subnet adjacency and credential signals, **not** observed network traffic — and Fenrir says so.
- Ranked by your **shared ACR criticality currency**, so “crown jewel” means what your business says it means.
- **Gated, reviewable writes** — Fenrir proposes tags; a human approves.

PROOF POINTS

In a reference environment of 268 assets, Fenrir mapped:

- **156** exploitable / weak-auth entry points.
- **67** crown-jewel targets ($ACR \geq 7$).
- **63 reachable attack paths** connecting them — a ranked worklist instead of a wall of findings.

Illustrative results from a demo lab — not a guarantee.

Reachability is inferred, not observed; signals depend on the relevant plugins being present, and blind or uncredentialed hosts are called out, not hidden.

WORKS BETTER WITH

Fenrir consumes **Laelaps** (KEV exploitability), **Janus** (credential weaknesses and identity pivots), **Anubis** (calibrated crown-jewel ACR), and **Charon/Mimir** (exploitable and EOL software) — the pack's correlation payoff in a single view.

WHO IT'S FOR

SOC and threat-led vulnerability teams; red and purple teams validating paths; CISOs who need exposure reporting that says “here's the chain, here's how we broke it.”

CALL TO ACTION

Ask Fenrir which chain reaches your crown jewels — then break it first.