



## MEET THE HOUNDS

# LAELAPS

## CISA KEV Tagging

*Laelaps Hound · Category: Threat & Exposure*

***“Every actively-exploited vuln, caught and tagged.”***

The hound fated to always catch its quarry — Laelaps never loses the scent of an actively-exploited vulnerability.

### THE PROBLEM

CISA's Known Exploited Vulnerabilities catalog is the single best “fix this now” signal in security — these aren't theoretical CVEs, they're being exploited in the wild today. But pulling KEV out of thousands of findings and grouping it into something a patch team can act on is tedious, manual, and easy to let slide.

### WHAT IT DOES

Laelaps instantly finds and tags every actively-exploited vulnerability in the estate. It works off the **CISA-KNOWN-EXPLOITED** cross-reference to tag all KEV assets at once, by catalog date, or **by month in one click** — and answers natural-language KEV hunts on top. The result is an exploitability backbone the entire pack keys off.

### KEY CAPABILITIES

- **KEV asset tagging** — every asset carrying an actively-exploited vulnerability is tagged from the CISA-KNOWN-EXPLOITED cross-reference in your own findings, putting the estate's most urgent population one filter away.
- **By-date and by-month roll-up tagging** — slice the KEV backlog by catalog date, or tag an entire month in one click. A quarter's remediation debt becomes three clean, schedulable batches.
- **Per-month distinct asset counts** — see exactly how many assets each month's wave touches before committing a patch window, so scoping happens before work starts, not during it.
- **Natural-language KEV queries** — ask “which assets picked up KEV findings in March?” in plain English; Laelaps translates to SQL and hunts. No query language required.
- **Gated writes to the tagging log** — every tag is proposed, human-approved, and recorded, so the exploitability backbone the rest of the pack depends on is also fully auditable.

### HOW IT WORKS

Laelaps runs against the local navi.db built from your Tenable data: the **vulns.xrefs** field (CISA-KNOWN-EXPLOITED) with catalog **dateAdded** parsing for date- and month-level grouping. All writes are proposed, human-approved, and logged.

### WHY IT'S DIFFERENT

- **Month-level bulk tagging** — no per-date clicking through the catalog; a quarter's KEV backlog becomes three tags.
- It's the **exploitability backbone of the pack** — the KEV signal Laelaps surfaces is what Fenrir, Pythia, and Heimdall prioritize against.
- **Evidence-based:** tags come from the KEV cross-reference in your actual findings, not a memorized CVE list.
- **Gated, reviewable writes** — Laelaps proposes tags; a human approves.

### PROOF POINTS

In a reference environment of 268 assets, Laelaps surfaced:

- **2,631 KEV findings** across the estate — instantly grouped and taggable by month instead of buried in the general finding count.

*Illustrative results from a demo lab — not a guarantee. KEV signals depend on scan coverage; blind or uncredentialed hosts are called out, not hidden.*

### WORKS BETTER WITH

Laelaps feeds the pack's heaviest hitters: **Fenrir** uses KEV as foothold evidence in attack paths, **Pythia** crosses it with AI assets, **Heimdall** folds it into crypto migration priority, and **On the Scent** rolls it into the executive exposure view.

### WHO IT'S FOR

Vulnerability management leads who need a defensible “patch this first” list, and patch prioritization owners who plan remediation in waves, not one CVE at a time.

### CALL TO ACTION

***Ask Laelaps what's actively exploited in your estate — it always catches its quarry.***