



MEET THE HOUNDS

MIMIR

Software Analyzer

Mimir Hound · Category: Discovery

“The memory of everything installed.”

Keeper of the well of wisdom — Mimir remembers every package, version, and lifecycle across the estate.

THE PROBLEM

Software inventory is fragmented — package data says one thing, CPE fingerprints say another, and neither tells the whole story. Teams can't see version sprawl, can't spot end-of-life exposure before it becomes unpatchable, and can't answer the simplest board question: “which product is dragging the most risk?”

WHAT IT DOES

Mimir merges the **software** and **CPE** inventories into one unified product map — with a **source delta** metric showing what each source missed — then enriches it with **endoflife.date** lifecycle data, flags **crown-jewel software**, builds a “behind latest” worklist, and ranks a **risk leaderboard** by KEV and critical exposure.

KEY CAPABILITIES

- **Merged product map with source badges** — reconciles the software and CPE inventories into one product list, each row badged with which source saw it — and a delta metric that quantifies what either source alone would have missed.
- **Version-sprawl, most-deployed, and single-install views** — instantly spot the product running fourteen versions, the package on half the estate, and the one-off install nobody owns.
- **EOL detection + tags** — **endoflife.date** lifecycle status lands directly on your inventory, before end-of-life becomes an audit finding or an unpatchable CVE.
- **Crown-jewel software flags** — role-based ACR suggestions raise the criticality of assets running the software the business depends on, feeding the pack's shared priority currency.
- **Risk leaderboard** — products ranked by the KEV and critical exposure they drag into the estate, so patching effort lands where risk actually concentrates.
- **Gated tags + CSV export** — every finding becomes a proposed, human-approved tag or an exportable worklist ready for the patch team.

HOW IT WORKS

Mimir runs against the local **navi.db** built from your Tenable data: the **software** and **cpes** tables reconciled into one map, enriched by a cached **endoflife.date** lifecycle feed, with the **vulns** table supplying risk. All writes are proposed, human-approved, and logged.

WHY IT'S DIFFERENT

- **Two inventory sources reconciled, with a delta metric** — Mimir finds what either source alone misses, and tells you how big the gap is.
- **Lifecycle + risk in one view** — EOL status and KEV exposure on the same row, instead of two spreadsheets and a prayer.
- **Feeds the pack:** its product map powers EOL tagging, AI discovery, and attack-path software evidence downstream.
- **Gated, reviewable writes** — Mimir proposes tags and ACR suggestions; a human approves.

PROOF POINTS

In a reference environment of 268 assets, Mimir reconciled:

- **~28,700 software rows** and **~43,700 CPE rows** into a single unified product map — with source badges showing exactly which inventory contributed each product.

Illustrative results from a demo lab — not a guarantee. Inventory depth depends on scan coverage; blind or uncredentialed hosts are called out, not hidden.

WORKS BETTER WITH

Mimir feeds **Charon** (EOL tagging from lifecycle data), **Anubis** (crown-jewel software → ACR calibration), **Pythia** (AI apps surfaced in CPEs), and **Fenrir** (exploitable software as attack-path evidence).

WHO IT'S FOR

Vulnerability management teams who need risk-ranked patching; IT asset managers reconciling inventories; patch teams that want one worklist instead of three exports.

CALL TO ACTION

Ask Mimir what's really installed — the well remembers everything.